

Business Entity Resolution

Amazon ML Challenge 2026 | Team Vulcans | 27 September 2026

Problem and design goal

For each deduplicated Source 1 business, predict all matching Source 2 and Source 3 IDs, including an empty list for singletons. Training has 2.21 million Source 1 records, 10.32 million targets, and 7.64 million labeled links. Test adds France with no French training labels. The metric is macro F0.5 over Source 1 groups, so false merges and missed singletons matter alongside retrieval recall.

1. Retrieve candidates with complementary routes

A local DuckDB index stores Unicode name and address tokens, target frequencies, and compact legal-suffix name keys. Each query uses rare name and address token postings within the same country. Separate name and address Jaro-Winkler top-100 quotas protect matches when either field is corrupted or absent. Compact and accent-folded name equality rescue punctuation and French accents. For India, an additional top-50 address-token-overlap channel helps cross-script names. Country is an open-set string, so France is processed without hard-coded exclusion.

2. Learned last-stage block, then final matching

A frozen 23-feature LightGBM ranks the broad union and passes at most 40 candidates per Source 1. A local word TF-IDF search adds up to 20 IDs per India query. For US queries with at most three initial matches, a second exact sparse TF-IDF address search adds up to five IDs. A character 3-4-gram name TF-IDF search adds up to ten IDs for US queries with at most two initial matches. The recorded final candidate TSV is the union actually scored. Top 40 retained 70,735/70,741 broad-reachable true links in development; the India quota added 1,346 reachable links. No pretrained weights or external data are used for these searches.

A separate 35-feature Unicode/transliteration LightGBM was trained on 5.25 million pairs mined from 190,086 training-owned queries. A 36-feature specialist replaces its score for targets with empty addresses, using target core-name frequency. India/other thresholds are 0.65/0.75; blank-address threshold is 0.80.

3. Resolve group and ownership conflicts

Rank selected links by model probability and retain at most 11 per Source 1, the largest true group in the full training labels. Assign a target predicted by multiple Source 1 records to its highest-scoring owner. Both passes preserve the exact candidate subset contract. A fixed address-number veto was rejected for this model because it reduced labeled F0.5.

Retrieval checks that changed the design

- The broad lexical union had 4.74 million development pairs and an oracle macro F0.5 ceiling of 0.96793. The learned top-40 block kept the ceiling at 0.96788.
- A full-index India address TF-IDF search found 1,870/3,070 missed true links among its top 100 address keys. Its bounded quota improved measured end-to-end F0.5 on development and separate frozen validation.
- The gated US top-five address route added 28 true links and zero false links on frozen validation; the same final matcher scored every added pair.
- The gated US top-ten character-name route added another seven true links and zero false links after the US address route on frozen validation.

Measured evidence and submission

All scores below are local macro F0.5; public and private leaderboard scores must be measured separately.

Held-out evaluation

Check	Scope	Result
Combined candidate oracle	22,133 development queries	0.96793 upper bound
Learned top-40 oracle	22,133 development queries	0.96788 upper bound
Original 23-feature baseline	Full development, 3 output passes	0.88331
Hard + blank specialist	Development, top 40, cap + owner	0.92322
Plus India address top 20	Development, cap + owner	0.93183
Hard + blank specialist	Frozen 1,994-query validation	0.91845
Plus India address top 20	Frozen 1,994-query validation	0.92752
Plus gated US address top 5	Frozen 1,994-query validation	0.92989
Plus gated US name top 10	Frozen 1,994-query validation	0.93118

What the measurements mean

- The candidate oracle assumes perfect scoring on retrieved pairs. It measures the retrieval ceiling, not an achieved model score.
- The learned top-40 block controls most candidate volume. India address top 20 raised frozen macro F0.5 by 0.00907; gated US address top 5 added 0.00237. Gated US character-name top 10 added 0.00129. Every quota was measured before full test inference.
- The general matcher compares original Unicode and local ASCII-transliterated name/address views, plus soft number agreement. The blank-address specialist uses locally computed core-name rarity. No French labels guided either model.
- France has no truth labels. In a 20,000-query test slice, generic-name collisions remained before capping. The top-11 cap was retained; no French F0.5 is claimed.

Reproducibility and fair play

The final ZIP contains the exact last-stage candidate TSV, matching TSV, runnable Python source, four locally trained model weights, pinned dependencies, and the full methodology. Both broad lexical and TF-IDF indexes are reproducible intermediates from supplied records. Every predicted ID belongs to the submitted candidate set. Streaming and official validators are run before submission. No supplied record text is sent to a remote service.

Team

Yash Kailas Doke, Harsh Jitendra Jain, Ayush Tiwari, and Vedant Kaulgekar.